

Temporal Difference Learning

- Sutton, R.S. and Barto, A.G., Reinforcement Learning: An Introduction, MIT Press, Cambridge, MA, 1998.
 - <http://webdocs.cs.ualberta.ca/~sutton/book/ebook/the-book.html>
 - Bible in this area.

Acknowledgement:

- Some of the slides in this chapter are partially modified from those by Hsin-Ti Tsai (蔡心迪) and Kun-Hao Yeh (葉騏豪).

Outline

- Reinforcement Learning
- Temporal Difference Learning
- Case Studies
 - 2048
 - Connect6

Reinforcement Learning

- A **computational approach** to learning from **interaction**
 - Explore designs for machines that are effective in
 - ▶ solving learning problems of scientific or economic interest,
 - ▶ evaluating the designs through mathematical analysis or computational experiments.
 - Focus on **goal-directed learning** from interaction, when compared with other approaches to machine learning.
 - The learner must discover which actions yield the most reward by trying them.
 - ▶ Two characteristics: most important distinguishing features of reinforcement learning.
 - **trial-and-error search**
 - **delayed reward**
 - different from **supervised learning**, like statistical pattern recognition, and artificial neural networks.




Successful Examples

- In AI, it has been used to defeat human champions at games of skill (Tesauro, 1994);
 - For Connect6/2048/Threes!, it has been used to reach the top levels.
 - For Go, it has been used in Monte-Carlo Tree Search.
- In robotics, to fly stunt maneuvers in robot-controlled helicopters (Abbeel et al., 2007).
- In neuroscience it is used to model the human brain (Schultz et al., 1997);
- In psychology to predict animal behavior (Sutton and Barto, 1990).
- In economics, it is used to understand the decisions of human investors (Choi et al., 2007), and to build automated trading systems (Nevmyvaka et al., 2006)
- In engineering, it has been used to allocate bandwidth to mobile phones and to manage complex power systems (Ernst et al., 2005).

Agent-Environment Interaction Framework

- Agent: The learner and decision-maker.
- Environment: The thing it interacts with, comprising everything outside the agent.
- State: whatever information is available to the agent.
- Reward: single numbers.



Agent-Environment Interaction

- The agent and environment interact at each of a sequence of discrete time steps, $t = 0, 1, 2, \dots$, or $t = 0, 1, 2, \dots T$ (T : the terminated time, if any.)
 - $S = \{s_0, s_1, s_2, \dots, s_t, \dots\}$.
 - ▶ s_t : some representation of the environment's state at time step t .
 - a_t : action at time step t ,
 - r_t : rewards at time step t ,
 - π_t : the agent's policy,
 - ▶ a mapping from states to probabilities of selecting each possible action
 - ▶ $\pi_t(s, a)$: the probability that $a_t = a$ if $s_t = s$.



Examples

- 2048-like games:
 - Make moves with rewards. Then, tiles are popped up randomly.
- Bioreactor:
 - Suppose reinforcement learning is being applied to determine moment-by-moment temperatures and stirring rates for a bioreactor
- Pick-and-Place Robot:
 - Consider using reinforcement learning to control the motion of a robot arm in a repetitive pick-and-place task.
- Recycling Robot
 - A mobile robot has the job of collecting empty soda cans in an office environment.
 - This agent has to decide whether the robot should
 - ▶ (1) actively search for a can for a certain period of time,
 - ▶ (2) remain stationary and wait for someone to bring it a can, or
 - ▶ (3) head back to its home base to recharge its battery.



Example: Recycling Robot


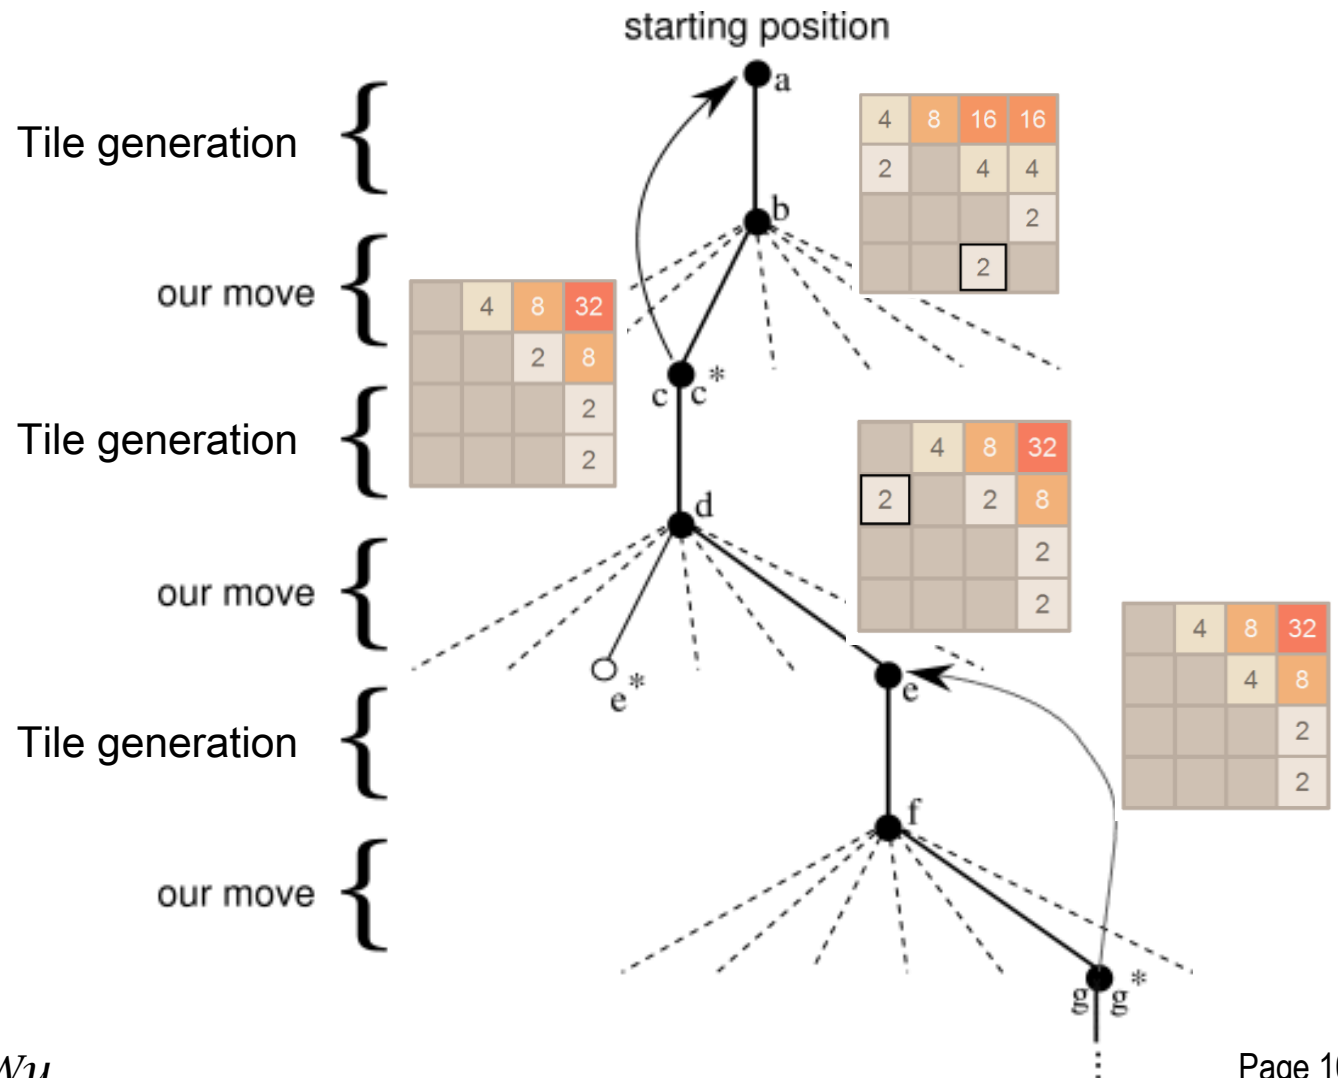


Figure 3.3: Transition graph for the recycling robot example.


$s = s_t$	$s' = s_{t+1}$	$a = a_t$	$\mathcal{P}_{ss'}^a$	$\mathcal{R}_{ss'}^a$
high	high	search	α	$\mathcal{R}^{\text{search}}$
high	low	search	$1 - \alpha$	$\mathcal{R}^{\text{search}}$
low	high	search	$1 - \beta$	-3
low	low	search	β	$\mathcal{R}^{\text{search}}$
high	high	wait	1	$\mathcal{R}^{\text{wait}}$
high	low	wait	0	$\mathcal{R}^{\text{wait}}$
low	high	wait	0	$\mathcal{R}^{\text{wait}}$
low	low	wait	1	$\mathcal{R}^{\text{wait}}$
low	high	recharge	1	0
low	low	recharge	0	0.



2048



Tic-Tac-Toe



Rewards

- Rewards: A way of formulating goal:
- Example: 2048
 - Straightforwardly set the earned scores to rewards.
- Example: recycling robot.
 - 0 for most of the time,
 - +1 for each can collected,
 - -3 in case of running out of electricity.
- Example: learning to play checkers/chess/Go,
 - +1 for winning,
 - -1 for losing, and
 - 0 for drawing and for all nonterminal positions.



Comments

- Critical:
 - **the rewards we set up truly indicate what we want accomplished.**
- Reward signals:
 - **A way of “what”** you want to achieve, **not “how”** you want to it achieved. (no prior knowledge about “how”)
 - Not the place to impart to the agent prior knowledge about how to achieve what we want it to do.
 - Example: for chess-playing, rewarded only for real winning, not for sub-goals, like taking pieces.



Goals

- Goal: Maximize the *expected return*.
- Returns: Total rewards of the episode
 - $R_t = r_{t+1} + r_{t+2} + r_{t+3} + \dots$
- Episodic tasks: (with a terminal state)
 - Example: 2048, chess, Go, etc.
 - $R_t = r_{t+1} + r_{t+2} + r_{t+3} + \dots + r_T$
 - ▶ T: a final time step. (s_T is a terminal state.)
 - ▶ Episode: $S^+ = \{s_{t+1}, s_{t+2}, s_{t+3}, \dots, s_T\}$.
Note that: $S = \{s_{t+1}, s_{t+2}, s_{t+3}, \dots, s_{T-1}\}$.
- Continuing tasks: $T=\infty$.
 - Example: Recycling Robot.
 - Problem: R_t could be infinite.
 - Solution: Add the concept of “discounting”.
 - Change it to “discounted return”:
 - $R_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots$
 - ▶ γ : discount rate. $0 \leq \gamma \leq 1$



Goal

- Rewards: A way of formulating goal:
 - Example: 2048
 - ▶ Straightforwardly set the earned scores to rewards.
- Goal: Maximize the *expected return*.
 - Returns: Total rewards of the episode
 - ▶ $R_t = r_{t+1} + r_{t+2} + r_{t+3} + \dots + r_T$
 - T: a final time step. (s_T is a terminal state.)
 - ▶ Example: 2048, chess, Go, etc.
- Episodic tasks: (with a terminal state)
 - Example: 2048, chess, Go, etc.
 - $R_t = r_{t+1} + r_{t+2} + r_{t+3} + \dots + r_T$
 - ▶ T: a final time step. (s_T is a terminal state.)
 - ▶ Episode: $S^+ = \{s_{t+1}, s_{t+2}, s_{t+3}, \dots, s_T\}$.
Note that: $S = \{s_{t+1}, s_{t+2}, s_{t+3}, \dots, s_{T-1}\}$.



Temporal Difference (TD) Learning

- Temporal Difference Learning
 - Is a kind of reinforcement learning
 - Is to adjust weights automatically
 - Was successfully applied to many games such as
 - ▶ Backgammon
 - ▶ Checkers
 - ▶ Chess
 - ▶ Shogi
 - ▶ Go
 - ▶ Chinese Chess
 - ▶ Connect6
 - ▶ 2048



Value Function

- Expected return or

Estimate how good it is for the agent to be in a given state

- $V(s)$: the estimated value of state s .
 - ▶ the expected return when starting in s thereafter.
 - ▶ also called the state-value function.
- $Q(s, a)$: the value of taking action a in state s under a policy π .
 - ▶ the expected return starting from s , taking the action a :
 - ▶ called the action-value function.
- Omit policy π . (See Reinforcement Learning)

TD Prediction

● Prediction:

- $V(s_t)$ is approximate to actual return R_t .
- Error: $\delta_t = R_t - V(s_t)$.
- Adjust: $V(s_t) = V(s_t) + \alpha\delta_t = V(s_t) + \alpha(R_t - V(s_t))$
 - ▶ α : a step-size parameter to control the learning rate.

● Problem:

- To get R_t , we must wait until the episode ends.
- Can we do that earlier?

TD(0)

- Change error:
 - From: $\delta_t = R_t - V(s_t)$
 - To: $\delta_t = (r_{t+1} + V(s_{t+1})) - V(s_t)$
 $= r_{t+1} + V(s_{t+1}) - V(s_t)$
 - ▶ For simplicity, $\gamma = 1$.
- Thus, change value function
 - From: $V(s_t) = V(s_t) + \alpha(R_t - V(s_t))$
 - To: $V(s_t) = V(s_t) + \alpha(r_{t+1} + V(s_{t+1}) - V(s_t))$
- This is called TD(0).

TD(λ)

- Change error:

- From: $\delta_t = R_t - V(s_t)$

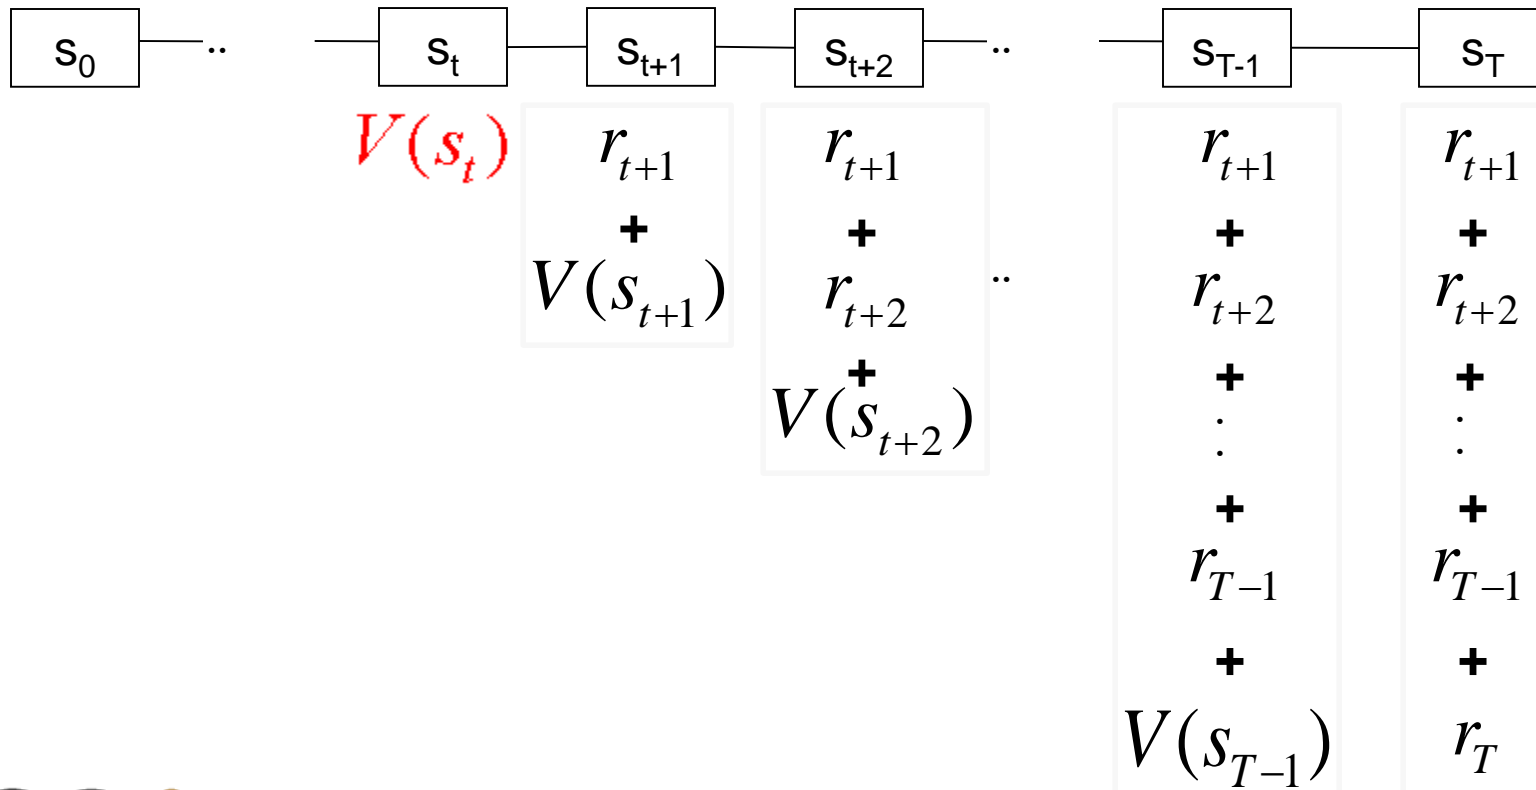
- To:

$$\delta_t = (1 - \lambda) \sum_{n=1}^{T-t-1} \lambda^{n-1} V(s_{t+n}) + \lambda^{T-t-1} V(s_T) - V(s_t)$$

- This is called TD(λ).


TD Learning

● TD(λ)




TD Learning

● TD(λ)



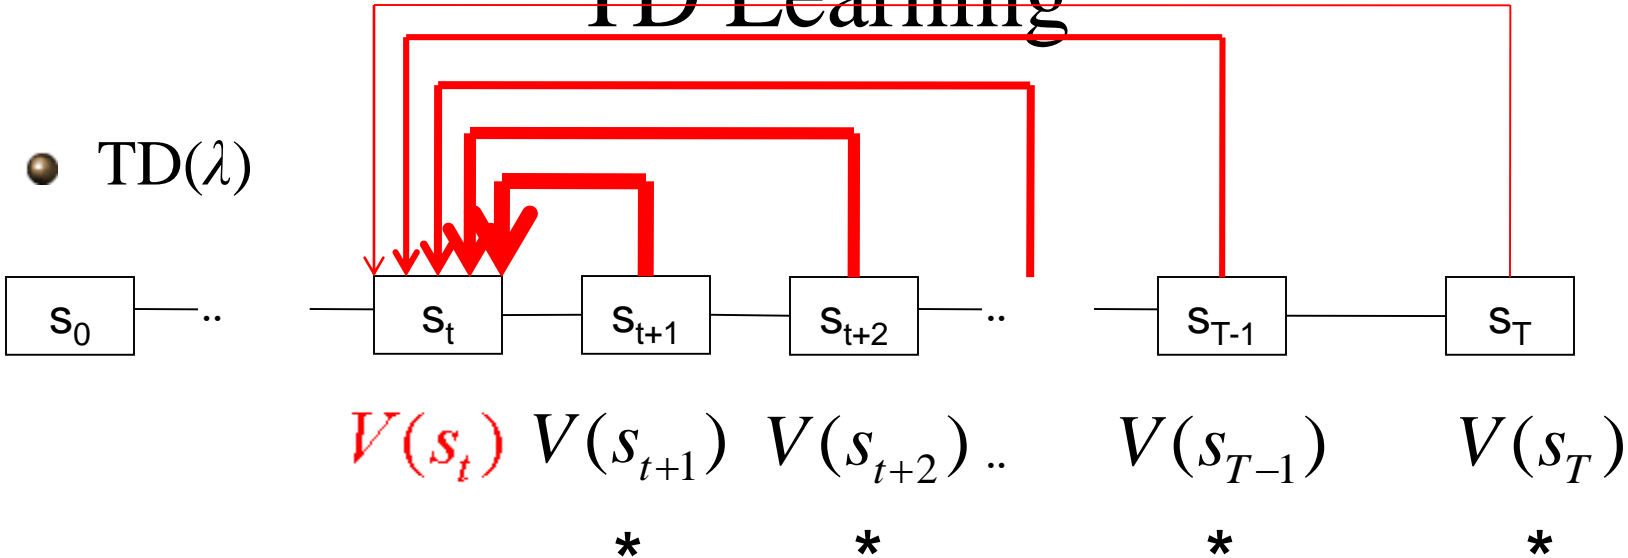
TD Learning

● TD(λ)



TD Learning

- TD(λ)



Proportion : $1-\lambda \quad (1-\lambda)\lambda \quad \dots \quad (1-\lambda)\lambda^{T-t-2} \quad \lambda^{T-t-1}$

$$\Sigma = 1$$

Given Weights





Figure 7.4: Weighting given in the λ -return to each of the n -step returns.




The Forward View



TD(1)


- Monte-Carlo tree search.



$$\Delta V(s_t) = \alpha[V(s_T) - V(s_t)]$$

$$V(s_t) = V(s_t) + \Delta V(s_t)$$

TD(0)




$$\Delta V(s_t) = \alpha[V(s_{t+1}) - V(s_t)]$$

$$V(s_t) = V(s_t) + \Delta V(s_t)$$

Case Study: 2048

- [Szubert and Jaskowski 2014]



N-Tuple Network

- Example: 8 4-tuple networks as shown.
 - Each cell has 16 different tiles
 - 16^4 features for this network.
 - ▶ But only one is on, others are 0.
 - ▶ So, we can use table lookup to find the feature weight.

64	⁰	8	4
128	¹		2
2	²		2
128	³		

0123	weight
0000	3.04
0001	-3.90
0002	-2.14
⋮	⋮
0010	5.89
⋮	⋮
0130	-2.01
⋮	⋮



Evaluation on Feature Weights

- General evaluation function:

- $V(s) = F(\varphi(s))$

- ▶ $\varphi(s)$: a vector of feature occurrences in s

- Linear evaluation function:

- $V(s) = \varphi(s) \cdot \theta$

- ▶ θ : a vector of feature weights



Evaluate a Position


- The value of a position is evaluated based on
 - linear combination of features, i.e.,
 - $V(s) = \varphi(s) \cdot \theta$
 - ▶ $\varphi(s)$: a vector of feature occurrences in s
 - ▶ θ : a vector of feature weights
- Features:
 - $\varphi(s)$: 8×16^4 features, $[0, 1, 0, \dots, 0, 0, 1, \dots, \dots, 1, 0, 0, \dots]$
 - ▶ All 0s, except for 8 ones.
 - One 1 every 16^4 features.
 - Let their indices be $f_1, f_2, f_3, f_4, f_5, f_6, f_7, f_8$.
- So, $V(s) = \varphi(s) \cdot \theta$

$$\begin{aligned}
 &= \sum_i^{8 \times 16^4} \varphi_i(s) \cdot \theta_i \\
 &= \sum_i^8 \varphi_{f_i}(s) \cdot \theta_{f_i} \\
 &= \sum_i^8 \theta_{f_i} \quad (\text{simply lookup table for } f_i)
 \end{aligned}$$



TD(0)

- To minimize the error.



- Error:

$$\delta_t = r_{t+1} + V(s_{t+1}) - V(s_t)$$

- ▶ So, $\Delta V(s_t) = \alpha \delta_t = \alpha (r_{t+1} + V(s_{t+1}) - V(s_t))$

- Adjustment

$$\Delta \theta = \Delta V(s_t) \frac{\varphi(s_t)}{\|\varphi(s_t)\|} = \alpha \delta_t \frac{\varphi(s_t)}{\|\varphi(s_t)\|}$$

- Since $\|\varphi(s_t)\|$ is constant in 2048, no need for normalization.



$$\Rightarrow \Delta \theta = \alpha' \delta_t \varphi(s_t),$$

- ▶ where $\alpha' = \frac{\alpha}{\|\varphi(s_t)\|}$



Three Methods of Evaluating Values

1. Evaluate actions: $Q(s, a)$.
 Select $\arg \max_a Q(s, a)$
 - Also called Q-learning
 - Problem: Too many features.
 - ▶ 4 times more!
 - ▶ This makes learning slower.




Three Methods of Evaluating Values

2. Evaluate states to play: $V(s_t)$.

Select $\arg_a \max(R(s_t, a) \sum_{s_{t+1}} P(s_t, a, s_{t+1}) V(s_{t+1}))$

– Problem: Higher time complexity.





Three Methods of Evaluating Values

3. Evaluate states after an action. $V(s'_t)$.

Select $\arg\max_a [R(s_t, a) + V(s'_t)]$

– These states are also called **afterstates**.

– **The best solution!!**



Afterstate Evaluation Function

```
1: function EVALUATE( $s, a$ )
2:    $s', r \leftarrow \text{COMPUTE AFTERSTATE}(s, a)$ 
3:   return  $r + V(s')$ 
4:
5: function LEARN EVALUATION( $s, a, r, s', s''$ )
6:    $a_{next} \leftarrow \arg \max_{a' \in A(s'')} \text{EVALUATE}(s'', a')$ 
7:    $s'_{next}, r_{next} \leftarrow \text{COMPUTE AFTERSTATE}(s'', a_{next})$ 
8:    $V(s') \leftarrow V(s') + \alpha(r_{next} + V(s'_{next}) - V(s'))$ 
```

Figure 6: The *afterstate evaluation function* and a dedicated variant of the TD(0) algorithm.




```
1: function PLAY GAME
2:    $score \leftarrow 0$ 
3:    $s \leftarrow \text{INITIALIZE GAME STATE}$ 
4:   while  $\neg \text{IS TERMINAL STATE}(s)$  do
5:      $a \leftarrow \arg \max_{a' \in A(s)} \text{EVALUATE}(s, a')$ 
6:      $r, s', s'' \leftarrow \text{MAKE MOVE}(s, a)$ 
7:     if LEARNING ENABLED then
8:       LEARN EVALUATION( $s, a, r, s', s''$ )
9:      $score \leftarrow score + r$ 
10:     $s \leftarrow s''$ 
11:   return  $score$ 
12:
13: function MAKE MOVE( $s, a$ )
14:    $s', r \leftarrow \text{COMPUTE AFTERSTATE}(s, a)$ 
15:    $s'' \leftarrow \text{ADD RANDOM TILE}(s')$ 
16:   return ( $r, s', s''$ )
```

Figure 3: A pseudocode of a game engine with moves selected according to the evaluation function. If learning is enabled, the evaluation function is adjusted after each move.





The N-Tuple Networks Used

- Use the following [Szubert and Jaskowski 2014]



- Ours:




Issues

- Features
 - Multi-stages
 - Monotonicity
 - Number of distinct tiles
 - Number of empty squares
 - Big tiles
 - Rotation/Mirroring
 - Sizes
- Step-size parameter: α .
 - Our experience: 0.0025,
 - ▶ A better version: 0.00025 after 1,000,000 learning games.
- Learning backwards.
- Bitboard
- Expectimax search

What If $V(s)$ is non-linear?

- Use gradient: $\nabla V(s_t)$
 - The Normal (法向量)
 - ▶ (等高線圖的梯度最大者.)
- Adjustment:
 - $\Delta\theta = \Delta V(s_t) \times \nabla V(s_t) = \alpha \delta_t \times \nabla V(s_t)$
- Example:
 - For linear: $\nabla V(s_t) = \frac{\varphi(s_t)}{\|\varphi(s_t)\|}$




Example

- $x = [x_1, x_2]$

$$f(x) = -4X_1^2 - X_2^2$$


- Find the max.
- Starting at (1,3)



Example

$$f'(x) = [-8X_1, -2X_2]$$


- From (1,3).
- Gradient: $[-8, -6]$
- Adjustment:
 - $\alpha = 0.1$
 - $(1,3) + \alpha(-8,-6)$
 $= (0.2, 2.4)$



Example

$$f'(x) = [-8X_1, -2X_2]$$

- From (0.2 , 2.4)
- Gradient: [-1.6 , -4.8]
- Adjustment:
 - $\alpha = 0.1$
 - $(0.2, 2.4) + \alpha(-1.6, -4.8)$
= (0.04 , 1.92)




Case Study: Connect6

- Connect6 is a kind of six-in-a-row game.
 - **Two** players, named Black and White.
 - ▶ Place **two** black and white stones, respectively.
 - Black plays first and places **one** stone initially.
 - ▶ A player wins
 - If the player gets **six or more** consecutive stones horizontally, vertically or diagonally.


Evaluate a Position

- Basically, we evaluate the value of a position based on the features of Connect6.
 - Usually use **linear combination** $V(s) = \varphi(s) \cdot \theta$,
 - ▶ $\varphi(s)$: a vector of feature occurrences in s ,
 - ▶ θ : a vector of feature weights



Evaluate a Position (An Illustration)


- For simplicity of discussion, use linear combination to evaluate the value of a position. E.g.,
 - Black: $1600 + 200*2 +$



T1: 1600
L3: 800
D3: 400
L2: 200
D2: 100
L1: 50

Evaluate a Position (An Illustration)


- For simplicity of discussion, use linear combination to evaluate the value of a position. E.g.,
 - Black: $1600 + 200*2 + 50*7 = 2350$



T1:	1600
L3:	800
D3:	400
L2:	200
D2:	100
L1:	50

Evaluate a Position (An Illustration)

- For simplicity of discussion, use linear combination to evaluate the value of a position. E.g.,
 - Black: $1600 + 200*2 + 50*7 = 2350$
 - White: $400 + 100*2 + 50*5 = 850$
 - Value for Black $= 2350 - 850 = 1500$




T1: 1600
 L3: 800
 D3: 400
 L2: 200
 D2: 100
 L1: 50



Goal

- For simplicity of discussion, use linear combination to evaluate the value of a position. E.g.,

- Black: $1600 + 200*2 + 50*7 = 2350$
- White: $400 + 100*2 + 50*5 = 850$
- Value for Black $= 2350 - 850 = 1500$




T1:	1600
L3:	800
D3:	400
L2:	200
D2:	100
L1:	50

We want to adjust these weights for accuracy.




TD Learning

- A sequence of positions.



TD(0)

- To minimize the error.



- Error:

$$\delta_t = V(s_{t+1}) - V(s_t)$$
$$\Delta V(s_t) = \alpha \delta_t = \alpha (V(s_{t+1}) - V(s_t))$$

- ▶ Note: reward is at the final state s_T

- Adjustment:

$$\Delta \theta = \Delta V(s_t) \frac{\varphi(s_t)}{\|\varphi(s_t)\|} = \alpha \delta_t \frac{\varphi(s_t)}{\|\varphi(s_t)\|}$$

- ▶ $\|\varphi(s_t)\|$ is for normalization.




Threat Space Search (TSS)

- We use TSS to **remove** the winning positions found by TSS. (to avoid updating from these positions)



TSS



Trained Weights

- The weights of some features after training.

Feature Weights	With TSS		Without TSS
W_{T2}	0.52982	Not high & close	1.73220
W_{T1}	0.51070		0.83796
W_{L3}	0.49358		0.73046
W_{D3}	0.27506		0.25531
W_{L2}	0.20028		0.07715

- T2: double threats (or live 4)
- T1: single threat (or dead 4)

- “without TSS” overweighs threats.



Discussion

- We successfully use TD(0) to improve the strength of NCTU6, a Connect6 champion program.
 - We got 58% win rate against the original NCTU6.
- We raise an important issue.
 - It is very important to remove the winning/losing positions found by TSS (or RZOP) in TD Learning.