Alpha-Beta Search

References:

- D. E. Knuth and R. W. Moore. An analysis of alpha-beta pruning. Artificial Intelligence, 6:293–326, 1975.
- J. Pearl. The solution for the branching factor of the alpha-beta pruning algorithm and its optimality. Communications of ACM, 25(8):559–564, 1982.

Acknowledgement:

• The slides of this chapter are modified from Prof. Hsu's teaching material, under the courtesy of Hsu.

http://www.iis.sinica.edu.tw/~tshsu/tcg2007/index.html

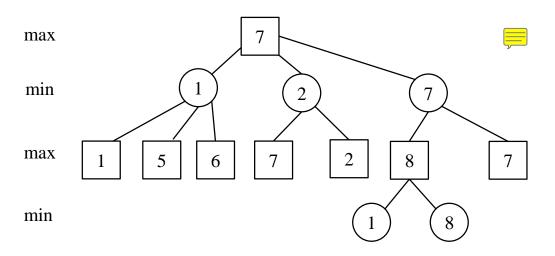


Introduction

- Alpha-beta pruning is the standard searching procedure used for 2person perfect-information zero sum games.
- Definitions:
 - A position *p*.
 - The value of a position p: f(p), is a numerical value computed from evaluating p.
 - ▶ Value is computed from the root player's point of view.
 - ▶ Positive values mean in favor of the root player.
 - Negative values mean in favor of the opponent.
 - Since it is a zero sum game, thus from the opponent's point of view, the value can be assigned, g(p) = -f(p).
 - A terminal position: a position whose value can be know.
 - ▶ Can be a position where win/loss/draw can be concluded.
 - ▶ Can be a position where some constraints are met.
 - A position p has d legal moves p_1, p_2, \ldots, p_d .



Mini-max formulation:



Mini-max formulation

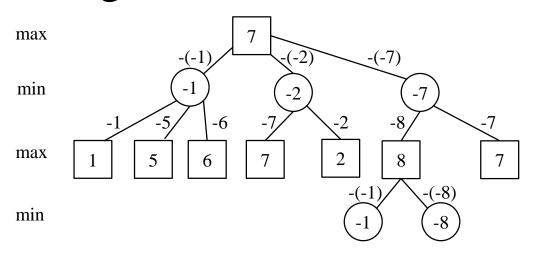
$$F(p) = \begin{cases} f(p) & \text{if } d = 0 \\ max\{G(p_1), \dots, G(p_d)\} & \text{if } d > 0 \end{cases}$$

$$G(p) = \begin{cases} g(p) & \text{if } d = 0 \\ min\{F(p_1), \dots, F(p_d)\} & \text{if } d > 0 \end{cases}$$

- An indirect recursive formula!
- Equivalent to AND-OR logic.



Nega-max formulation



- Nega-max formulation:
 - Let F(p) be the greatest possible value achievable from position p against the optimal defensive strategy.

$$F(p) = \begin{cases} f(p) & \text{if } d = 0\\ max\{-F(p_1), \dots, -F(p_d)\} & \text{if } d > 0 \end{cases}$$

Equivalent to NOR or NAND logic.



Algorithm: Nega-max

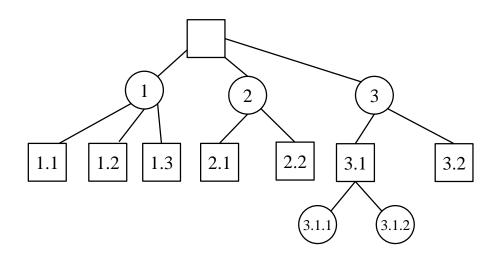
• Algorithm F(position p)determine the successor positions p_1, \ldots, p_d if d = 0, then return f(p) else begin $m := -\infty$ for i := 1 to d do $t = -F(p_i)$ if t > m then m = t \rightleftharpoons end

• A brute-force method to try all possibilities!



return m

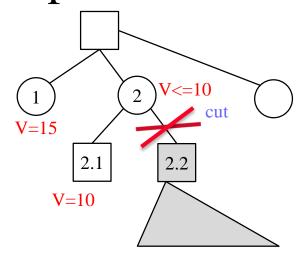
Representation



- From the root, number the node in a search tree by a sequence of integers a.b.c.d · · ·
 - Meaning from the root, you first take the ath branch, then the bth branch, and then the cth branch, and then the dth branch $\cdot \cdot \cdot$
 - The root is specified as an empty sequence.
 - This is called "Dewey decimal system".



Alpha Cut-off

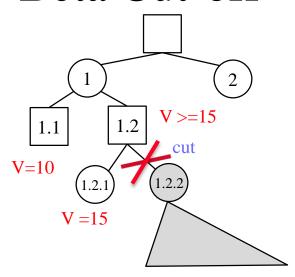


On a max node

- Assume you have finished exploring the first branch 1 and obtain the best value from that branch as bound.
- You now search the second branch 2.
- Assume branch at 2.1 returns a value of < bound.
 - ▶ Then no need to evaluate 2.2, 2.3, ..., at all.
- Since the best possible value for 2 is <= bound (10),
 - the root should choose the first branch as the current best solution.



Beta Cut-off



On a min node

- Assume you have finished exploring the first branch 1.1 at 1 and obtain the best value from that branch as bound (10).
- You now search the second branch 1.2.
- Assume branch at 1.2.1 returns a value (15) of > bound.
 - ▶ Then no need to evaluate 1.2.2, 1.2.3, ..., at all.
- Since the best possible value for 1.2 is > bound,
 - choose the first branch as the current best solution.

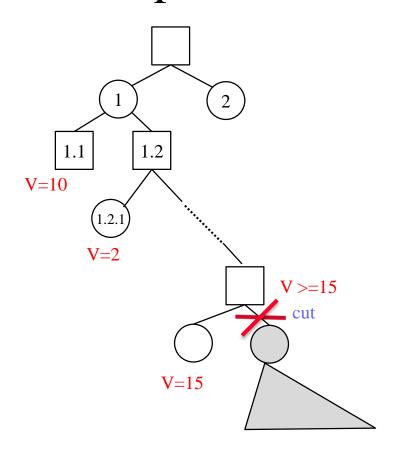


Deep Cut-off

- Shallow (alpha) cut-off (See slides before)
 - For a min node u, the branch of its elder brother produces a lower bound V_l for its parent.
 - The first branch of u produces an upper bound V_u for v.
 - If $V_l > V_u$, then there is no need to evaluate the second branch of u.
- Deep (alpha) cut-off:
 - Def: For a node u in a tree and a positive integer g, Ancestor(g, u) is the direct ancestor of u by tracing the parent's link g times.
 - When the lower bound V_l is produced at and propagated from u's great grand parent, i.e., Ancestor(3,u), or any Ancestor(2i+1,u).
- Similar properties for beta cut-off.



Deep Cut-off





Illustration

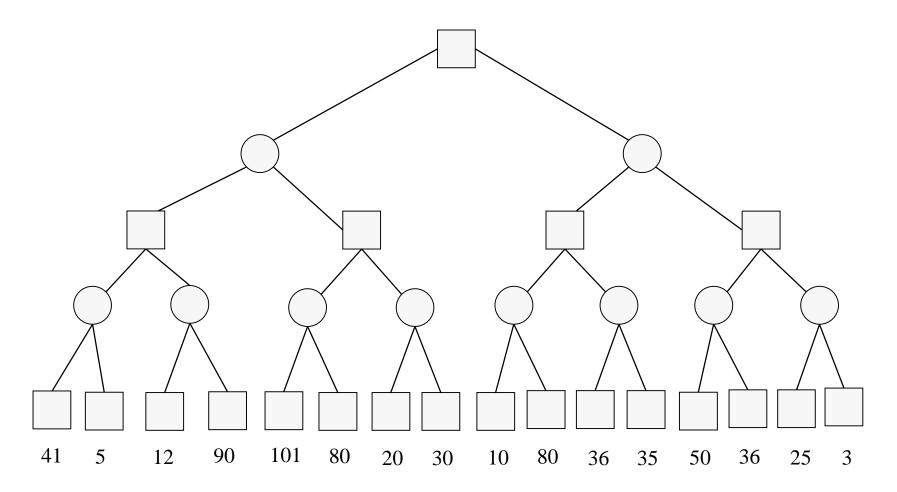




Illustration of Alpha-Beta Search







Illustration of Alpha-Beta Search





Alpha-beta Pruning Algorithm: Mini-Max

Algorithm F2 '(position p, integer alpha, integer beta) determine the successor positions p_1, \ldots, p_d if d = 0, then return f(p) else begin m := alphafor i := 1 to d do $t := G2'(p_i, m, beta)$ if t > m then m := tif $m \ge$ beta then return(m) //cutoff end; return *m* Algorithm G2 '(position p, integer alpha, integer beta) determine the successor positions p_1, \ldots, p_d if d = 0, then return g(p) else begin m := betafor i := 1 to d do $t := F2'(p_i, alpha, m)$ if t < m then m := tif $m \leq \text{alpha then return}(m) //\text{cutoff}$



Alpha-beta pruning algorithm: Nega-max

• Algorithm F2(position p, integer alpha, integer beta)

```
determine the successor positions p_1, \ldots, p_d if d = 0, then return f(p) else begin  m := alpha  for i := 1 to d do  t := -F2(p_i, -beta, -m)  if t > m then m := t if m \ge beta then return(m) end return m
```



Properties and Comments

• Properties:



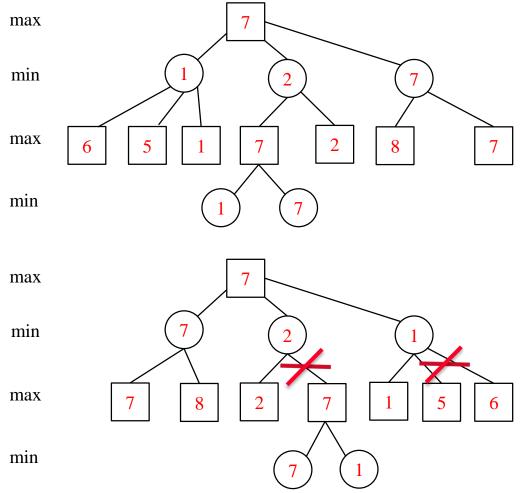
- alpha < beta
- $F2(p, alpha, beta) \leq alpha, \text{ if } F(p) \leq alpha$
- F2(p, alpha, beta) = F(p), if alpha < F(p) < beta
- $F2(p, alpha, beta) \ge beta$, if $F(p) \ge beta$
- $-F2(p,-\infty,+\infty)=F(p)$

Comments:

- -F2(p, alpha, beta): find the best possible value according to a mini-max formula for the position p with the constraints that
 - ▶ If F(p) is less than the lower bound *alpha*, then F2(p, alpha, beta) returns a terminal positions value that is less than *alpha*.
 - ▶ If F(p) is more than the upper bound *beta*, then F2(p, alpha, beta) returns a terminal positions value that is more than *beta*.
- The meanings of *alpha* and *beta* during searching:
 - For a max node: the current best value is at least *alpha*.
 - For a min node: the current best value is at most *beta*.



Examples with Different Ordering





Analysis of a Possible Best Case

- Q: In the best case, what branches are cut?
- Definitions:
 - A path in a search tree is a sequence of numbers indicating the branches selected in each level using the Dewey decimal system.
 - A position is denoted as a path $a_1.a_2. \cdot \cdot \cdot .a_l$ from the root.
 - A position $a_1.a_2. \cdot \cdot \cdot .a_l$ is critical if
 - $a_i = 1$ for all even values of i or
 - $a_i = 1$ for all odd values of i
 - Examples: 2.1.4.1.2, 1.3.1.5.1.2, 1.1.1.2.1.1.1.3 and 1.1 are critical
 - ► Examples: 1.2.1.1.2 is not critical
 - A perfect-ordering tree:

$$F(a_1.\cdots.a_\ell) = \begin{cases} f(a_1.\cdots.a_\ell) & \text{if } a_1.\cdots.a_\ell \text{ is a terminal} \\ -F(a_1.\cdots.a_\ell.1) & \text{otherwise} \end{cases}$$

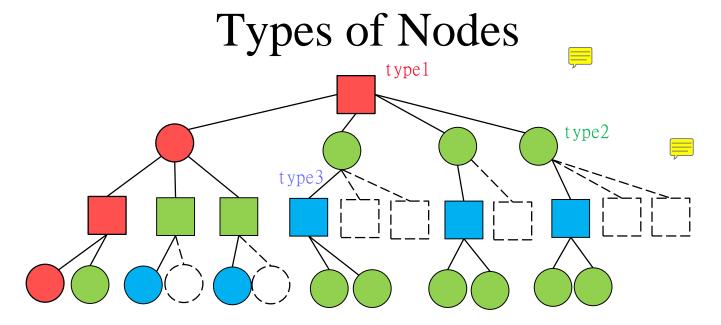
▶ The first successor of every non-terminal position gives the best possible value.



Critical Positions

- The minimum positions needed to derive the value.
- Theorem 1: F2 examines precisely the critical positions of a perfect-ordering tree.





- Classification of critical positions $a_1.a_2. \cdot \cdot \cdot .a_1$:
 - type 1: all the a_i are 1;
 - type 2: if a_j is its first entry such that $a_j > 1$, for all odd k, $a_{j+k} = 1$;
 - ▶ All the children of type 1 nodes, except for the first.
 - ▶ All the children of type 3 nodes.
 - type 3: the first child of type 2 nodes.



Proof Sketch for Theorem 1

- Properties (invariants)
 - A type 1 position p is examined by calling $F2(p, -\infty, \infty)$
 - p's first successor p_1 is type $1 \rightarrow F(p_1) = -F(p_1) \neq \pm \infty$
 - ▶ p's other successors p_2, \ldots, p_d are of type 2 → p_i , i > 1, are examined by calling $F2(p_2, -\infty, F(p_2))$
 - A type 2 position p is examined by calling $F2(p, -\infty, beta)$ where $-\infty < beta \le F(p)$
 - p's first successor p_1 is type $3 \rightarrow F(p) = -F(p_1)$
 - p's other successors p_2, \ldots, p_d are not examined
 - A type 3 position *p* is examined by calling F2(p, alpha, ∞) where ∞ > $alpha \ge F(p)$
 - p's successors p_1, \ldots, p_d are of type 2
 - they are examined by calling $F2(p_i, -\infty, -alpha)$
- Using an induction argument to prove all and also only critical positions are examined.



Analysis: Best Case

- Corollary 1: Assume each position has exactly d successors
 - The alpha-beta procedure examines exactly $d^{\lceil l/2 \rceil} + d^{\lfloor l/2 \rfloor} 1$ positions on level l.

• Proof:

- There are $d^{\lceil l/2 \rceil}$ sequences of the form $a_1.a_2. \cdot \cdot \cdot .a_l$ with $1 \le a_i \le d$ for all i such that $a_i = 1$ for all odd values of i.
- There are $d^{\lfloor l/2 \rfloor}$ sequences of the form $a_1.a_2. \cdot \cdot \cdot .a_l$ with $1 \leq a_i \leq d$ for all i such that $a_i = 1$ for all even values of i.
- We substrate 1 for the sequence $1.1. \cdot \cdot \cdot \cdot 1.1$ which are counted twice.



Analysis: Average Case

- Assumptions: Let a random game tree be generated in such a way that
 - each position on level j has probability q_i of being nonterminal
 - has an average of d_i successors
- Properties of the above random game tree
 - Expected number of positions on level l is $d_0.d_1. \cdot \cdot \cdot .d_{l-1}$.
 - Expected number of positions on level *l* examined by an alpha-beta procedure assumed the random game tree is perfectly ordered is

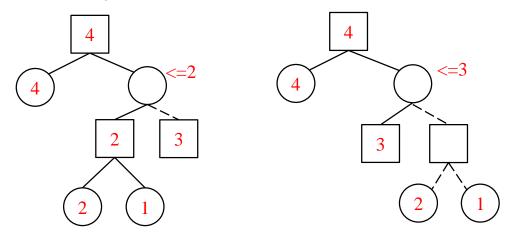
$$d_0q_1d_2q_3 \cdot \cdot \cdot d_{l-2}q_{l-1} + q_0d_1q_2d_3 \cdot \cdot \cdot q_{l-2}d_{l-1} + q_0q_1...q_{l-1}$$
 if l is even; (skipped for odd)

- Proof sketch:
 - If x is the expected number of positions of a certain type on level j, then xd_j is the expected number of successors of these positions, and xq_j is the expected number of "number 1" successors.
 - The above numbers equal to those of Corollary 1 when $q_j = 1$ and $d_j = d$ for $0 \le j < l$.



Perfect Ordering is Not Always Best

- Intuitively, we may "think" alpha-beta pruning would be most effective when a game tree is perfectly ordered.
 - That is, when the first successor of every position is the best possible move.
 - This is not always the case!



• Truly optimum order of game trees traversal is not obvious.



Theorem 2

- Theorem 2: Alpha-beta pruning is optimum in the following sense:
 - Given any game tree and any algorithm which computes the value of the root position, there is a way to permute the tree
 - by reordering successor positions if necessary;
 - so that every terminal position examined by the alpha-beta method under this permutation is examined by the given algorithm.
 - Furthermore if the value of the root is not 1 or −1, the alpha-beta procedure examines precisely the positions which are critical under this permutation.



Questions

- What is a good move ordering?
 - It may not be good to search the best possible move first.
 - It is best to cut off a branch with more nodes.
- How about the case when the tree is not uniform?
- What is the effect of using iterative-deepening alpha-beta cut off?
- How about the case for searching a game graph instead of a game tree?
 - Can some nodes be visited more than once?



History

- McCarthy thought of the method during the Dartmouth Summer Research Conference on Artificial Intelligence in 1956.
 - No formal specification of the algorithm was given at that time.
 - McCarthy's remarks at that conference led to the use of alpha-beta pruning in game-playing programs of the late 1950s.
 - [15] McCarthy, J. Personal communication, December 1, 1973.
- Samuel has stated that the idea was present in his checker-playing programs.
 - but he did not allude to it in his classic article [21] because he felt that the other aspects of his program were more significant.
 - [21] Samuel, A. L. Some studies in machine learning using the game of checkers. IBMJ. Res. and Develop. 3 (1959).
- The first published discussion of a method for game tree pruning appeared in Newell, Shaw and Simon's description [16] of their early chess program.
 - However, they illustrate only the "one-sided" technique used in procedure F1 above, so it is not clear whether they made use of "deep cutoffs".
- McCarthy coined the name "alpha-beta" when he first wrote a LISp program embodying the technique.
 - Hart and Edwards, who wrote a memorandum [10] on the subject in 1961.
- The first published account of alpha-beta pruning actually appeared in Russia, quite independently of the American work, in 1963.
 - [4] Brudno, A, L. Bounds and valuations for shortening the scanning of variations. Problemy Kibernet. 10 (1963), 141-150 (in Russian).
- Excellent presentations of the method appear in the textbooks by Nilsson [18, Section 4] and Slagle [23, pp. 16-24],
 - [18] Nilsson, N, J. Problem-Solving Methods in Artificial Intelligence. McGraw-Hill, New York. 1971.
- This paper is still valuable even when
 - alpha-beta pruning has been in use for more than 15 years

